

```
import pandas as pd
import numpy as np
```

```
df=pd.read_csv('Customer_Data.csv')
print(df.head())
```

	Customer ID	Age	Gender	Annual Income	Spending Score
0	1	62	Male	56564	86
1	2	65	Male	25393	73
2	3	18	Male	143520	39
3	4	21	Male	64711	43
4	5	21	Male	147014	98

```
print("Shape of the dataset:",df.shape)
print("Column names:",df.columns)
```

```
Shape of the dataset: (1000, 5)
Column names: Index(['Customer ID', 'Age', 'Gender', 'Annual Income',
'Spending Score'], dtype='object')
```

```
missing=df.isnull().sum()
print(missing)
```

```
Customer ID      0
Age              0
Gender           0
Annual Income    0
Spending Score   0
dtype: int64
```

```
print(df.describe())
```

	Customer ID	Age	Annual Income	Spending Score
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	43.267000	84856.809000	47.859000
std	288.819436	15.242311	38393.323903	28.606038
min	1.000000	18.000000	20359.000000	1.000000
25%	250.750000	30.000000	51173.750000	23.000000
50%	500.500000	43.000000	84462.000000	47.000000
75%	750.250000	56.000000	118356.250000	73.000000
max	1000.000000	69.000000	149870.000000	99.000000

```
!pip install matplotlib
!pip install seaborn
import matplotlib.pyplot as plt
import seaborn as sns
```

```
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
```

Requirement already satisfied: cycloper>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: numpy>=1.20 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.25.2)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (24.0)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.16.0)
Requirement already satisfied: seaborn in
/usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (1.25.2)
Requirement already satisfied: pandas>=1.2 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (2.0.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (1.2.1)
Requirement already satisfied: cycloper>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (24.0)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in

```

/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4-
>seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn)
(2023.4)
Requirement already satisfied: tzdata>=2022.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn)
(2024.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)

```

```


sns.set(style="whitegrid")
plt.figure(figsize=(15,5))
#numerical variables visualization
plt.subplot(1,3,1)
sns.histplot(df['Age'],bins=10,kde=True,color='skyblue')
plt.title('Distribution of Age')

plt.subplot(1,3,2)
sns.histplot(df['Annual Income'],bins=10,kde=True,color='salmon')
plt.title('Distribution of Annual Income')

plt.subplot(1,3,3)
sns.histplot(df['Spending Score'],bins=10,kde=True,color='green')
plt.title('Distribution of Spending Score')

Text(0.5, 1.0, 'Distribution of Spending Score')


```



```

plt.figure(figsize=(8,5))
sns.countplot(data=df,x='Gender')
plt.title('Distribution of Gender')
plt.show()

```



```
!pip install scikit-learn
from sklearn.preprocessing import StandardScaler, LabelEncoder

Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.3.2 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)

label=LabelEncoder()
df['Gender']=label.fit_transform(df['Gender'])
print(df.head())
```

	Customer ID	Age	Gender	Annual Income	Spending Score
0	1	62	1	56564	86
1	2	65	1	25393	73
2	3	18	1	143520	39
3	4	21	1	64711	43
4	5	21	1	147014	98

```

scaler=StandardScaler()
num=['Age','Annual Income','Spending Score']
df[num]=scaler.fit_transform(df[num])
print(df.head())

```

	Customer ID	Age	Gender	Annual Income	Spending Score
0	1	1.229628	1	-0.737289	1.333987
1	2	1.426547	1	-1.549581	0.879310
2	3	-1.658518	1	1.528717	-0.309845
3	4	-1.461599	1	-0.524984	-0.169944
4	5	-1.461599	1	1.619768	1.753689

```

from sklearn.model_selection import train_test_split

```

```

X= df.drop(columns=['Customer ID'])
y=df['Customer ID']
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.2,random_state=42)
print(X_train.shape)
print(X_test.shape)

```

```

(800, 4)
(200, 4)

```

```

from sklearn.cluster import KMeans
kmeans=KMeans(random_state=42)
inertia=[]
for k in range(1,11):
    kmeans=KMeans(n_clusters=k,random_state=42)
    kmeans.fit(X_train)
    inertia.append(kmeans.inertia_)

```

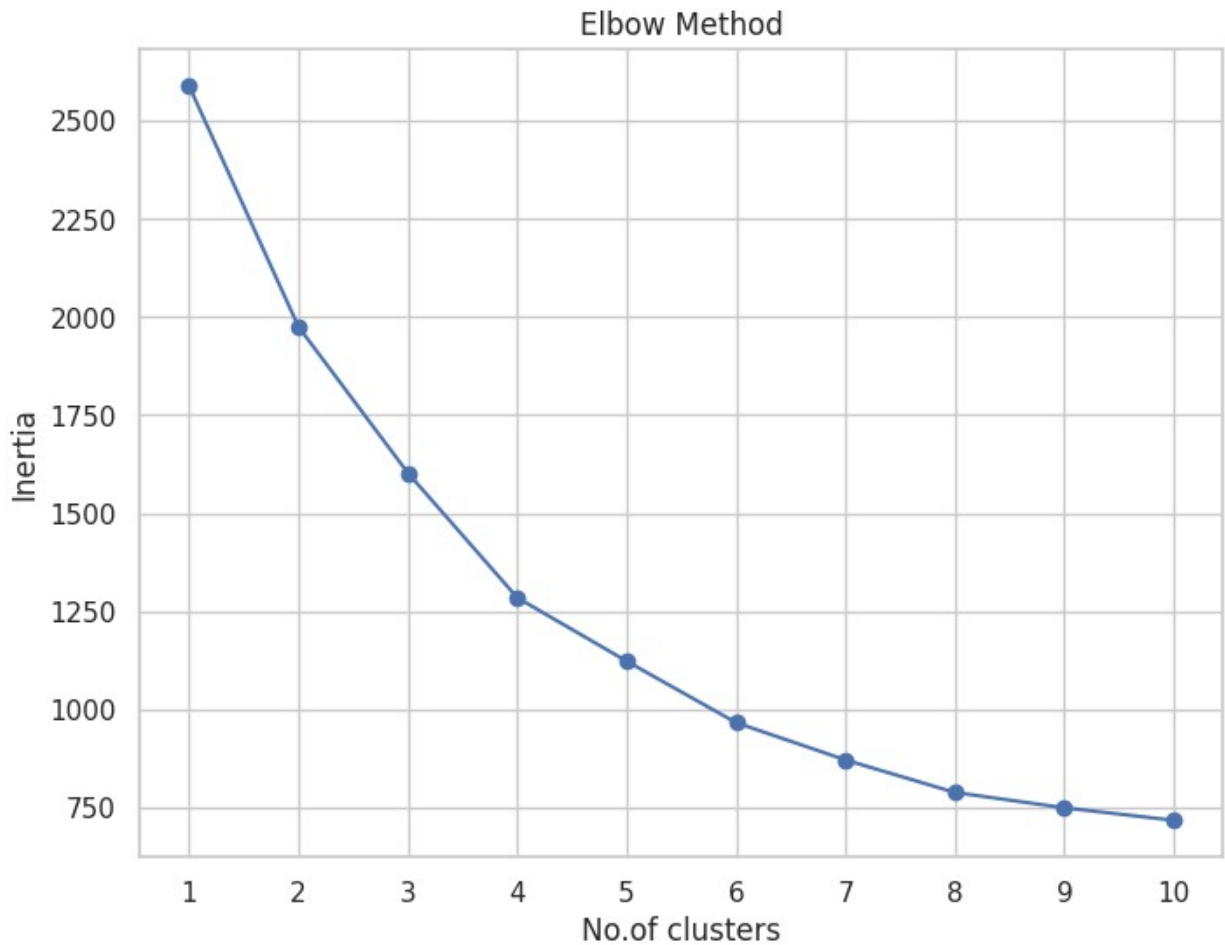
```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning

```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870
: FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning)

plt.figure(figsize=(8,6))
plt.plot(range(1,11),inertia,marker='o')
plt.title('Elbow Method')
plt.xlabel('No.of clusters')
plt.ylabel('Inertia')
plt.xticks(np.arange(1,11,1))
plt.show()
```

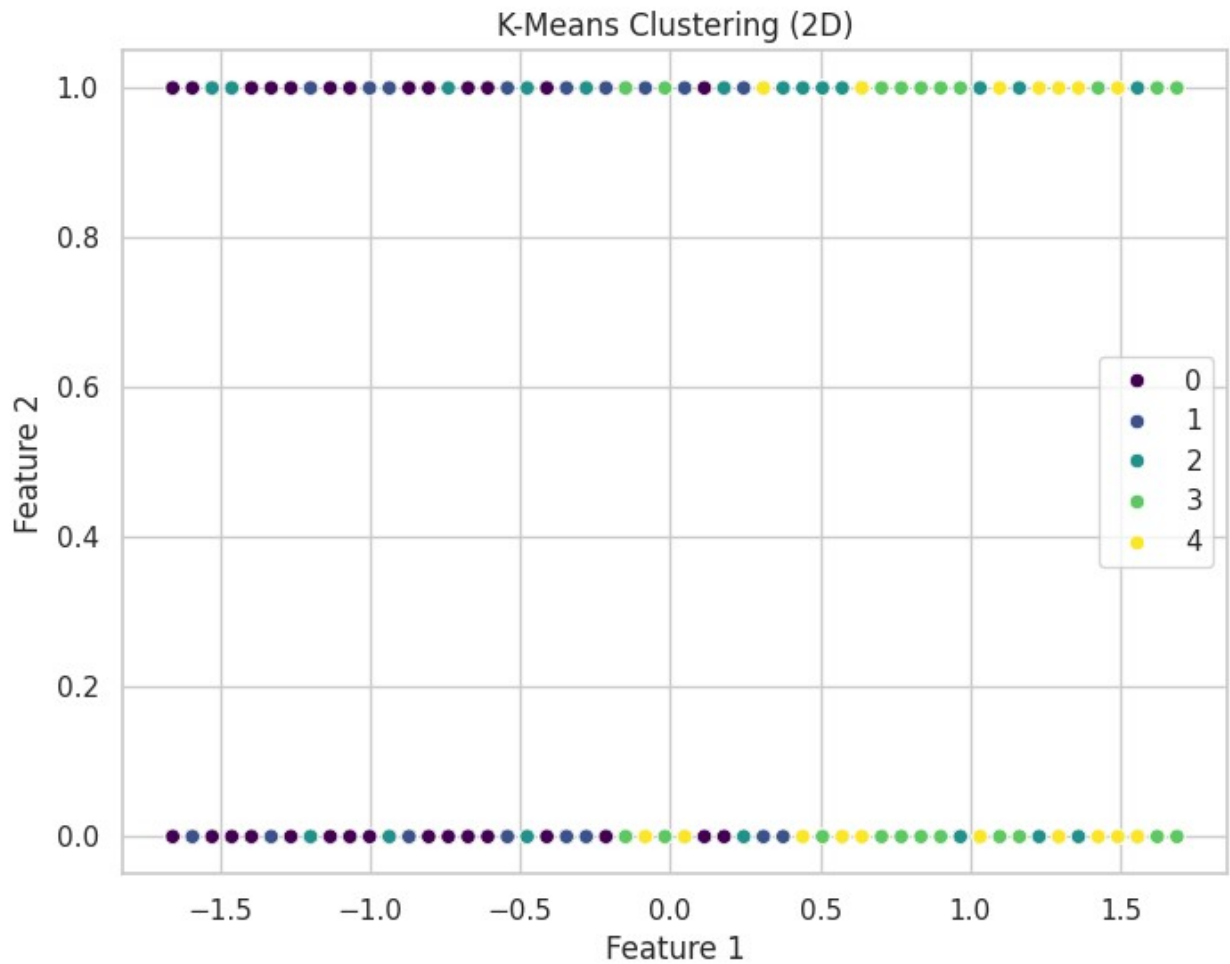


```
k=5
kmeans=KMeans(n_clusters=k,random_state=42)
kmeans.fit(X_train)

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
  warnings.warn(

KMeans(n_clusters=5, random_state=42)

plt.figure(figsize=(8,6))
sns.scatterplot(x=X_train.iloc[:,0],y=X_train.iloc[:,1],hue=kmeans.labels_,palette='viridis',legend='full')
plt.title('K-Means Clustering (2D)')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```



```
from sklearn.metrics import silhouette_score,davies_bouldin_score
silhouette=silhouette_score(X_train,kmeans.labels_)
print(silhouette)

0.23938393558789123

davies=davies_bouldin_score(X_train,kmeans.labels_)
print(davies)

1.2370478276957144

"""
Discussion of strengths and limitations of K-means clustering
Strengths:
- K-means is computationally efficient and scales well to large
datasets.
- It is simple to implement and easy to understand.
- K-means can produce tight clusters if the data is well-separated
and spherical in shape.

Limitations:
```


- K-means requires the number of clusters (k) to be specified in advance, which can be challenging in practice.
- It assumes that clusters are isotropic and have similar densities, which may not always hold true.
- K-means is sensitive to the initial choice of cluster centers and may converge to suboptimal solutions.

```
{"type": "string"}
```

"""

Insights and Recommendations

Based on the clustering results, we have identified distinct customer segments. Let's provide actionable insights and marketing strategies tailored to each segment:

Cluster 0: High Spending Customers

- Offer premium loyalty programs or rewards to encourage repeat purchases.
- Introduce exclusive products or services targeting this segment.
- Personalize marketing communications to enhance engagement.

Cluster 1: Moderate Spending Customers

- Provide targeted discounts or promotions to incentivize higher spending.
- Implement referral programs to increase customer acquisition.
- Improve product visibility through targeted advertising campaigns.

Cluster 2: Low Spending Customers

- Offer introductory discounts or bundle deals to attract this segment.
- Focus on improving the value proposition of products or services.
- Gather feedback to understand barriers to spending and address customer concerns.

Cluster 3: High Income, Low Spending Customers

- Create personalized experiences to entice this segment to spend more.
- Showcase the luxury aspects of products or services to appeal to their higher income.
- Offer exclusive perks or benefits to increase loyalty.

Cluster 4: Low Income, High Spending Customers

- Provide budget-friendly options or payment plans to accommodate their spending habits.
- Focus on building brand loyalty through exceptional customer service.
- Implement targeted marketing campaigns emphasizing affordability and value.

The company can use these insights to tailor marketing strategies and improve customer engagement:

- Personalize marketing messages and offers based on the characteristics of each customer segment.*
- Utilize customer segmentation for targeted advertising and promotional campaigns.*
- Continuously analyze customer feedback and interactions to refine marketing strategies and offerings.*

"""

```
{"type": "string"}
```